

# Non Linear SVM - Email Spam Classifier

October 22, 2020

## 1 Non-Linear SVM - Email Spam Classifier

We'll build a non-linear SVM classifier to classify emails and compare the performance with the linear SVM model.

To reiterate, the performance of the linear model was as follows: - accuracy 0.93 - precision 0.92 - recall 0.89

```
In [1]: import pandas as pd
import numpy as np
from sklearn.svm import SVC
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import GridSearchCV
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import scale
```

### Loading Data

```
In [2]: email_rec = pd.read_csv("Spam.csv", sep = ',')
```

### 1.1 Data Preparation

```
In [3]: # splitting into X and y
X = email_rec.drop("spam", axis = 1)
y = email_rec.spam.values.astype(int)
```

```
In [4]: # scaling the features
X_scaled = scale(X)
```

```
# train test split
```

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size = 0.3, rand
```

C:\Anaconda3\lib\site-packages\ipykernel\_launcher.py:2: DataConversionWarning: Data with input

## 1.2 Model Building

```
In [5]: # using rbf kernel, C=1, default value of gamma
```

```
model = SVC(C = 1, kernel='rbf')
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
```

```
C:\Anaconda3\lib\site-packages\sklearn\svm\base.py:196: FutureWarning: The default value of gamma
"avoid this warning.", FutureWarning)
```

## 1.3 Model Evaluation Metrics

```
In [6]: # confusion matrix
```

```
confusion_matrix(y_true=y_test, y_pred=y_pred)
```

```
Out[6]: array([[811,  38],
               [ 61, 471]], dtype=int64)
```

```
In [7]: # accuracy
```

```
print("accuracy", metrics.accuracy_score(y_test, y_pred))
```

```
# precision
```

```
print("precision", metrics.precision_score(y_test, y_pred))
```

```
# recall/sensitivity
```

```
print("recall", metrics.recall_score(y_test, y_pred))
```

```
accuracy 0.9283128167994207
```

```
precision 0.925343811394892
```

```
recall 0.8853383458646616
```

## 1.4 Hyperparameter Tuning

Now, we have multiple hyperparameters to optimise - - The choice of kernel (linear, rbf etc.) - C - gamma

We'll use the GridSearchCV() method to tune the hyperparameters.

## 1.5 Grid Search to Find Optimal Hyperparameters

Let's first use the RBF kernel to find the optimal C and gamma (we can consider the kernel as a hyperparameter as well, though training the model will take an exorbitant amount of time).

```
In [8]: # creating a KFold object with 5 splits
```

```
folds = KFold(n_splits = 5, shuffle = True, random_state = 4)
```

```
# specify range of hyperparameters
```

```

# Set the parameters by cross-validation
hyper_params = [ {'gamma': [1e-2, 1e-3, 1e-4],
                  'C': [1, 10, 100, 1000]}]

# specify model
model = SVC(kernel="rbf")

# set up GridSearchCV()
model_cv = GridSearchCV(estimator = model,
                        param_grid = hyper_params,
                        scoring= 'accuracy',
                        cv = folds,
                        verbose = 1,
                        return_train_score=True)

# fit the model
model_cv.fit(X_train, y_train)

```

Fitting 5 folds for each of 12 candidates, totalling 60 fits

```

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 60 out of 60 | elapsed: 46.5s finished

```

```

Out[8]: GridSearchCV(cv=KFold(n_splits=5, random_state=4, shuffle=True),
                    error_score='raise-deprecating',
                    estimator=SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
                                decision_function_shape='ovr', degree=3, gamma='auto_deprecated',
                                kernel='rbf', max_iter=-1, probability=False, random_state=None,
                                shrinking=True, tol=0.001, verbose=False),
                    fit_params=None, iid='warn', n_jobs=None,
                    param_grid=[{'gamma': [0.01, 0.001, 0.0001], 'C': [1, 10, 100, 1000]}],
                    pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                    scoring='accuracy', verbose=1)

```

```

In [9]: # cv results
cv_results = pd.DataFrame(model_cv.cv_results_)
cv_results

```

```

Out[9]:
   mean_fit_time  std_fit_time  mean_score_time  std_score_time  param_C  \
0      0.426284    0.020828      0.060368      0.003716         1
1      0.474971    0.022318      0.093137      0.009366         1
2      0.876207    0.103492      0.146991      0.002362         1
3      0.296357    0.018458      0.046820      0.003405        10
4      0.314254    0.010734      0.054949      0.004604        10
5      0.453887    0.010305      0.083359      0.002161        10
6      0.405114    0.097553      0.039683      0.007592       100

```

7	0.390674	0.050723	0.049930	0.006212	100
8	0.368195	0.017466	0.059163	0.001172	100
9	0.545270	0.038658	0.041219	0.009998	1000
10	0.540438	0.045874	0.035787	0.002707	1000
11	0.419198	0.083070	0.048977	0.010300	1000

	param_gamma	params	split0_test_score \
0	0.01	{'C': 1, 'gamma': 0.01}	0.917702
1	0.001	{'C': 1, 'gamma': 0.001}	0.886646
2	0.0001	{'C': 1, 'gamma': 0.0001}	0.770186
3	0.01	{'C': 10, 'gamma': 0.01}	0.909938
4	0.001	{'C': 10, 'gamma': 0.001}	0.917702
5	0.0001	{'C': 10, 'gamma': 0.0001}	0.883540
6	0.01	{'C': 100, 'gamma': 0.01}	0.913043
7	0.001	{'C': 100, 'gamma': 0.001}	0.923913
8	0.0001	{'C': 100, 'gamma': 0.0001}	0.919255
9	0.01	{'C': 1000, 'gamma': 0.01}	0.908385
10	0.001	{'C': 1000, 'gamma': 0.001}	0.919255
11	0.0001	{'C': 1000, 'gamma': 0.0001}	0.920807

	split1_test_score	split2_test_score	...	mean_test_score \
0	0.939441	0.922360	...	0.929814
1	0.919255	0.899068	...	0.904037
2	0.802795	0.791925	...	0.786025
3	0.944099	0.934783	...	0.933230
4	0.934783	0.916149	...	0.928261
5	0.914596	0.899068	...	0.902174
6	0.937888	0.934783	...	0.931677
7	0.940994	0.925466	...	0.933851
8	0.934783	0.917702	...	0.927019
9	0.922360	0.920807	...	0.918323
10	0.944099	0.930124	...	0.933851
11	0.936335	0.925466	...	0.929193

	std_test_score	rank_test_score	split0_train_score	split1_train_score \
0	0.008528	5	0.943323	0.940994
1	0.013080	10	0.910326	0.903339
2	0.015322	12	0.789208	0.779503
3	0.012266	3	0.966227	0.966615
4	0.009491	7	0.937112	0.932453
5	0.013749	11	0.909938	0.902174
6	0.010159	4	0.982531	0.979814
7	0.008482	1	0.950311	0.949534
8	0.007349	8	0.934006	0.931289
9	0.005607	9	0.993789	0.992624
10	0.009033	1	0.966615	0.966227
11	0.005777	6	0.940606	0.940994

	split2_train_score	split3_train_score	split4_train_score	\
0	0.945264	0.937112	0.939829	
1	0.908773	0.906056	0.904115	
2	0.785326	0.791925	0.788820	
3	0.967003	0.961568	0.962345	
4	0.936335	0.935171	0.931289	
5	0.908773	0.905280	0.902562	
6	0.982531	0.982143	0.982531	
7	0.948758	0.945652	0.939829	
8	0.934006	0.930901	0.929348	
9	0.992624	0.993012	0.992236	
10	0.966615	0.963121	0.966227	
11	0.940994	0.937112	0.937500	

	mean_train_score	std_train_score
0	0.941304	0.002814
1	0.906522	0.002672
2	0.786957	0.004277
3	0.964752	0.002308
4	0.934472	0.002242
5	0.905745	0.003158
6	0.981910	0.001059
7	0.946817	0.003835
8	0.931910	0.001831
9	0.992857	0.000527
10	0.965761	0.001331
11	0.939441	0.001753

[12 rows x 22 columns]

```
In [10]: # converting C to numeric type for plotting on x-axis
cv_results['param_C'] = cv_results['param_C'].astype('int')

# # plotting
plt.figure(figsize=(16,6))

# subplot 1/3
plt.subplot(131)
gamma_01 = cv_results[cv_results['param_gamma']==0.01]

plt.plot(gamma_01["param_C"], gamma_01["mean_test_score"])
plt.plot(gamma_01["param_C"], gamma_01["mean_train_score"])
plt.xlabel('C')
plt.ylabel('Accuracy')
plt.title("Gamma=0.01")
plt.ylim([0.80, 1])
plt.legend(['test accuracy', 'train accuracy'], loc='upper left')
plt.xscale('log')
```

```


# subplot 2/3
plt.subplot(132)
gamma_001 = cv_results[cv_results['param_gamma']==0.001]

plt.plot(gamma_001["param_C"], gamma_001["mean_test_score"])
plt.plot(gamma_001["param_C"], gamma_001["mean_train_score"])
plt.xlabel('C')
plt.ylabel('Accuracy')
plt.title("Gamma=0.001")
plt.ylim([0.80, 1])
plt.legend(['test accuracy', 'train accuracy'], loc='upper left')
plt.xscale('log')

# subplot 3/3
plt.subplot(133)
gamma_0001 = cv_results[cv_results['param_gamma']==0.0001]

plt.plot(gamma_0001["param_C"], gamma_0001["mean_test_score"])
plt.plot(gamma_0001["param_C"], gamma_0001["mean_train_score"])
plt.xlabel('C')
plt.ylabel('Accuracy')
plt.title("Gamma=0.0001")
plt.ylim([0.80, 1])
plt.legend(['test accuracy', 'train accuracy'], loc='upper left')
plt.xscale('log')

```



This plot reveals some interesting insights: - **High values of gamma** lead to **overfitting** (especially at high values of C); note that the training accuracy at gamma=0.01 and C=1000 reaches almost 99% - The **training score increases with higher gamma**, though the **test scores are comparable** (at sufficiently high cost, i.e.  $C > 10$ ) - The least amount of overfitting (i.e. difference between train and test accuracy) occurs at low gamma, i.e. a quite *simple non-linear model*

```
In [11]: # printing the optimal accuracy score and hyperparameters
best_score = model_cv.best_score_
best_hyperparams = model_cv.best_params_
```

```
print("The best test score is {0} corresponding to hyperparameters {1}".format(best_s
```

The best test score is 0.9338509316770186 corresponding to hyperparameters {'C': 100, 'gamma':

Though sklearn suggests the optimal scores mentioned above (gamma=0.001, C=100), one could argue that it is better to choose a simpler, more non-linear model with gamma=0.0001. This is because the optimal values mentioned here are calculated based on the average test accuracy (but not considering subjective parameters such as model complexity).

We can achieve comparable average test accuracy (~92.5%) with gamma=0.0001 as well, though we'll have to increase the cost C for that. So to achieve high accuracy, there's a trade-off between: - High gamma (i.e. high non-linearity) and average value of C - Low gamma (i.e. less non-linearity) and high value of C

We argue that the model will be simpler if it has as less non-linearity as possible, so we choose gamma=0.0001 and a high C=100.

### 1.5.1 Building and Evaluating the Final Model

Let's now build and evaluate the final model, i.e. the model with highest test accuracy.

```
In [13]: # specify optimal hyperparameters
best_params = {"C": 100, "gamma": 0.0001, "kernel": "rbf"}

# model
model = SVC(C=100, gamma=0.0001, kernel="rbf")

model.fit(X_train, y_train)
y_pred = model.predict(X_test)

# metrics
print(metrics.confusion_matrix(y_test, y_pred), "\n")
print("accuracy", metrics.accuracy_score(y_test, y_pred))
print("precision", metrics.precision_score(y_test, y_pred))
print("sensitivity/recall", metrics.recall_score(y_test, y_pred))
```

```
[[810  39]
 [ 60 472]]
```

```
accuracy 0.9283128167994207
precision 0.923679060665362
sensitivity/recall 0.8872180451127819
```

## 1.6 Conclusion

The accuracy achieved using a non-linear kernel is comparable to that of a linear one. Thus, it turns out that for this problem, **you do not really need a non-linear kernel**.